

# Classification Takeaway Assignment

Import necessary libraries

```
In [104... import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCAN
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, f1_score, precision_score, re
from sklearn.metrics import silhouette_score, davies_bouldin_score, calin

from imblearn.over_sampling import SMOTE
```

Load the dataset

```
In [105... df = pd.read_csv('Dataset of Diabetes .csv')
```

## Exploratory Data Analysis

```
In [106... df.head()
```

```
Out[106...
   ID  No_Pation  Gender  AGE  Urea  Cr  HbA1c  Chol  TG  HDL  LDL  VLDL
0  502      17975      F   50   4.7  46    4.9   4.2  0.9   2.4   1.4   0.5
1  735      34221      M   26   4.5  62    4.9   3.7  1.4   1.1   2.1   0.6
2  420      47975      F   50   4.7  46    4.9   4.2  0.9   2.4   1.4   0.5
3  680      87656      F   50   4.7  46    4.9   4.2  0.9   2.4   1.4   0.5
4  504      34223      M   33   7.1  46    4.9   4.9  1.0   0.8   2.0   0.4
```

```
In [107... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   ID          1000 non-null   int64
 1   No_Pation   1000 non-null   int64
 2   Gender      1000 non-null   object
 3   AGE         1000 non-null   int64
 4   Urea        1000 non-null   float64
 5   Cr          1000 non-null   int64
 6   HbA1c       1000 non-null   float64
 7   Chol        1000 non-null   float64
 8   TG          1000 non-null   float64
 9   HDL         1000 non-null   float64
10   LDL         1000 non-null   float64
11   VLDL        1000 non-null   float64
12   BMI         1000 non-null   float64
13   CLASS       1000 non-null   object
dtypes: float64(8), int64(4), object(2)
memory usage: 109.5+ KB
```

Check for null values

```
In [108... print(df.isnull().sum())
```

```
ID          0
No_Pation   0
Gender      0
AGE         0
Urea        0
Cr          0
HbA1c       0
Chol        0
TG          0
HDL         0
LDL         0
VLDL        0
BMI         0
CLASS       0
dtype: int64
```

```
In [109... print(df['CLASS'].value_counts())
```

```
CLASS
Y      840
N      102
P       53
Y         4
N          1
Name: count, dtype: int64
```

```
In [110... df['CLASS'] = df['CLASS'].str.strip() #making sure there are only 3 class
```

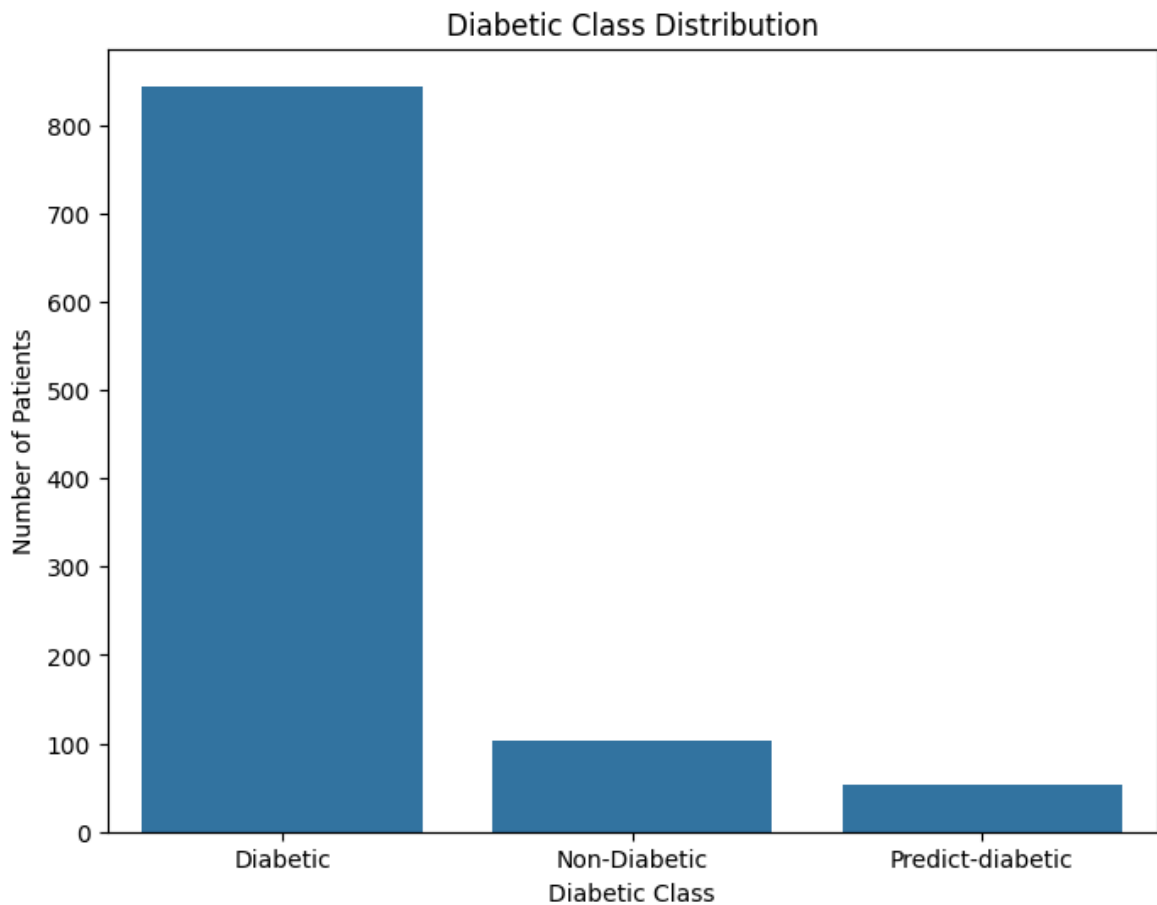
```
In [111... print(df['CLASS'].unique())
```

```
['N' 'P' 'Y']
```

```
In [112... class_dist = df['CLASS'].value_counts()
class_dist = class_dist.rename({"Y":"Diabetic", "P":"Predict-diabetic", "
```

```
plt.figure(figsize=(8, 6))
ax = sns.barplot(x=class_dist.index, y=class_dist.values)
plt.title("Diabetic Class Distribution")
plt.ylabel("Number of Patients")
plt.xlabel("Diabetic Class")
```

Out[112... Text(0.5, 0, 'Diabetic Class')



```
In [113... print(df['Gender'].value_counts())
```

```
Gender
M    565
F    434
f      1
Name: count, dtype: int64
```

```
In [114... df['Gender'] = df['Gender'].str.upper()
```

```
In [115... print(df['Gender'].unique())
```

```
['F' 'M']
```

```
In [116... le = LabelEncoder()
```

```
# Fit and transform the categorical data
df['Gender'] = le.fit_transform(df['Gender'])
df['CLASS'] = le.fit_transform(df['CLASS'])
```

```
In [117... df.describe()
```

Out [117...

	ID	No_Pation	Gender	AGE	Urea	
<b>count</b>	1000.000000	1.000000e+03	1000.000000	1000.000000	1000.000000	1000.0
<b>mean</b>	340.500000	2.705514e+05	0.565000	53.528000	5.124743	68.9
<b>std</b>	240.397673	3.380758e+06	0.496005	8.799241	2.935165	59.9
<b>min</b>	1.000000	1.230000e+02	0.000000	20.000000	0.500000	6.0
<b>25%</b>	125.750000	2.406375e+04	0.000000	51.000000	3.700000	48.0
<b>50%</b>	300.500000	3.439550e+04	1.000000	55.000000	4.600000	60.0
<b>75%</b>	550.250000	4.538425e+04	1.000000	59.000000	5.700000	73.0
<b>max</b>	800.000000	7.543566e+07	1.000000	79.000000	38.900000	800.0

In [118... `df.drop(columns=['ID', 'No_Pation'])`

Out [118...


	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS
<b>0</b>	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0
<b>1</b>	1	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	0
<b>2</b>	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0
<b>3</b>	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0
<b>4</b>	1	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	0
<b>...</b>	...	...	...	...	...	...	...	...	...	...	...	...
<b>995</b>	1	71	11.0	97	7.0	7.5	1.7	1.2	1.8	0.6	30.0	2
<b>996</b>	1	31	3.0	60	12.3	4.1	2.2	0.7	2.4	15.4	37.2	2
<b>997</b>	1	30	7.1	81	6.7	4.1	1.1	1.2	2.4	8.1	27.4	2
<b>998</b>	1	38	5.8	59	6.7	5.3	2.0	1.6	2.9	14.0	40.5	2
<b>999</b>	1	54	5.0	67	6.9	3.8	1.7	1.1	3.0	0.7	33.0	2

1000 rows × 12 columns

```
In [119... attribute_columns = ['AGE', 'Urea', 'Cr', 'HbA1c', 'Chol', 'TG', 'HDL', 'LDL', 'VLDL', 'BMI', 'CLASS']
corr_matrix = df[attribute_columns].corr()

plt.figure(figsize=(8, 8))
sns.heatmap(corr_matrix, annot=True, cmap='crest')
plt.title('Correlation Heatmap')
```


Out [119... `Text(0.5, 1.0, 'Correlation Heatmap')`



```
In [120...] diabetic_patients = df.loc[df["CLASS"] == 2]
non_diabetic_patients = df.loc[df["CLASS"] == 0]
predict_diabetic_patients = df.loc[df["CLASS"] == 1]
```

```
In [121...] age_count = diabetic_patients['AGE'].value_counts()
age_count = age_count.sort_index()

plt.figure(figsize=(15, 10))
bax = sns.barplot(x=age_count.index, y=age_count.values)
plt.title("Diabetic Patient's Age distribution")
plt.ylabel("Number of Patients")
plt.xlabel("Age")
plt.show()
```



## Model Building

```
In [122... X = df.drop('CLASS', axis=1) # Features
           y = df['CLASS']    # Target variable
```

```
In [123... X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## Classification Algorithms:

### K Nearest Neighbour (KNN)

```
In [124... training_accuracy = []
test_accuracy = []
training_f1 = []
test_f1 = []

neighbors_settings = range(2,20)

for n_neighbors in neighbors_settings:
    print(f'working on neighbors {n_neighbors}')
    knn = KNeighborsClassifier(n_neighbors=n_neighbors)
    knn.fit(X_train, y_train)

    y_train_pred=knn.predict(X_train)
    y_pred=knn.predict(X_test)

    training_accuracy.append(accuracy_score(y_train,y_train_pred))
```

```
test_accuracy.append(accuracy_score(y_test, y_pred))

training_f1.append(f1_score(y_train,y_train_pred, average='weighted'))
test_f1.append(f1_score(y_test, y_pred, average='weighted'))
```

```
working on neighbors 2
working on neighbors 3
working on neighbors 4
working on neighbors 5
working on neighbors 6
working on neighbors 7
working on neighbors 8
working on neighbors 9
working on neighbors 10
working on neighbors 11
working on neighbors 12
working on neighbors 13
working on neighbors 14
working on neighbors 15
working on neighbors 16
working on neighbors 17
working on neighbors 18
working on neighbors 19
```

```
In [125... results_df = pd.DataFrame({
    'k': neighbors_settings,
    'Training Accuracy': training_accuracy,
    'Test Accuracy': test_accuracy,
    'Training F1 Score': training_f1,
    'Test F1 Score': test_f1
})

print(results_df)

best_k_accuracy = results_df.loc[results_df['Test Accuracy'].idxmax()]
print(f"\nBest k based on Test Accuracy: {best_k_accuracy['k']}")
print(f"Test Accuracy: {best_k_accuracy['Test Accuracy']}")
print(f"Training Accuracy: {best_k_accuracy['Training Accuracy']}")

best_k_f1 = results_df.loc[results_df['Test F1 Score'].idxmax()]
print(f"\nBest k based on Test F1 Score: {best_k_f1['k']}")
print(f"Test F1 Score: {best_k_f1['Test F1 Score']}")
print(f"Training F1 Score: {best_k_f1['Training F1 Score']}")
```

	k	Training Accuracy	Test Accuracy	Training F1 Score	Test F1 Score
0	2	0.90000	0.795	0.909521	0.815808
1	3	0.88750	0.860	0.881025	0.856264
2	4	0.87750	0.850	0.874846	0.853246
3	5	0.87750	0.875	0.868952	0.870631
4	6	0.87625	0.865	0.868044	0.862955
5	7	0.87375	0.880	0.860605	0.874531
6	8	0.86375	0.875	0.855023	0.871945
7	9	0.86750	0.885	0.854716	0.879932
8	10	0.86125	0.870	0.851648	0.867626
9	11	0.86125	0.870	0.848364	0.865702
10	12	0.85875	0.850	0.848273	0.848139
11	13	0.85500	0.850	0.841479	0.837011
12	14	0.85875	0.850	0.843524	0.834067
13	15	0.86000	0.855	0.844496	0.841294
14	16	0.86125	0.845	0.845471	0.830618
15	17	0.85750	0.850	0.842732	0.840303
16	18	0.85375	0.845	0.839795	0.833434
17	19	0.85625	0.855	0.841749	0.843980

Best k based on Test Accuracy: 9.0

Test Accuracy: 0.885

Training Accuracy: 0.8675

Best k based on Test F1 Score: 9.0

Test F1 Score: 0.8799319727891156

Training F1 Score: 0.8547156862473493

```
In [126... fig = plt.figure(figsize=(14,10))
fig.subplots_adjust(hspace=0.5, wspace=0.3)

fig.add_subplot(2,2,1)
plt.plot(neighbors_settings, training_accuracy, label='training accuracy')
plt.plot(neighbors_settings, test_accuracy, label='test accuracy')
plt.ylabel('Accuracy',size=20)
plt.xticks(size=15)
plt.yticks(size=15)
plt.xlabel('n_neighbors',size=20)
plt.title('Accuracy Score',size=20)
plt.legend(['Training','Testing'],frameon=False, ncol=2);

fig.add_subplot(2,2,2)
plt.plot(neighbors_settings, training_f1)
plt.plot(neighbors_settings, test_f1)
plt.xticks(size=15)
plt.yticks(size=15)
plt.ylabel('F1 Score',size=20)
plt.xlabel('n_neighbors',size=20)
plt.title('F1-Score',size=20)
plt.legend(['Training','Testing'],frameon=False, ncol=2);
```






```
In [127... knn = KNeighborsClassifier(n_neighbors=9).fit(X_train, y_train)
y_pred = knn.predict(X_test)

print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.57	0.57	0.57	21
1	0.80	0.40	0.53	10
2	0.93	0.95	0.94	169
accuracy			0.89	200
macro avg	0.77	0.64	0.68	200
weighted avg	0.88	0.89	0.88	200

```
In [128... def normalized_confusion_matrix(y_test, conf_mat, model):
_ , counts = np.unique(y_test,return_counts=True)
conf_mat = conf_mat/counts
plt.figure(figsize=(6,5))
ax=sns.heatmap(conf_mat,fmt='.2f',annot=True,annot_kws={'size':20},lw
plt.title(f'Confusion Matrix ({model})',size=22)
plt.xticks(size=20)
plt.yticks(size=20)
```

```
In [129... conf_mat = confusion_matrix(y_test,y_pred)
normalized_confusion_matrix(y_test,conf_mat, 'kNN')
```




### Support Vector Machine (SVM)

```
In [130...] svc = SVC(C=1000, gamma='auto').fit(X_train_scaled, y_train)
y_pred=svc.predict(X_test_scaled)

print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	21
1	0.86	0.60	0.71	10
2	0.98	0.99	0.99	169
accuracy			0.96	200
macro avg	0.89	0.83	0.85	200
weighted avg	0.96	0.96	0.96	200

```
In [131...] conf_mat = confusion_matrix(y_test,y_pred)
normalized_confusion_matrix(y_test,conf_mat,'SVM')
```



Decision Tree

```
In [132... max_depth=range(1,20)
training_accuracy = []
test_accuracy = []
training_f1 = []
test_f1 = []

for depth in max_depth :
    tree = DecisionTreeClassifier(random_state=0, max_depth=depth, min_sa
    y_train_pred = tree.predict(X_train)
    y_pred = tree.predict(X_test)

    training_accuracy.append(accuracy_score(y_train,y_train_pred))
    test_accuracy.append(accuracy_score(y_test, y_pred))

    training_f1.append(f1_score(y_train,y_train_pred,average='weighted'))
    test_f1.append(f1_score(y_test, y_pred,average='weighted'))
```

```
In [133... best_max_depth = pd.DataFrame({
    "Max Depth": max_depth,
    "Train Accuracy": training_accuracy,
    "Test Accuracy": test_accuracy,
    "Train F1": training_f1,
    "Test F1": test_f1
})


print(best_max_depth.to_string(index=False))
```

Max Depth	Train Accuracy	Test Accuracy	Train F1	Test F1
1	0.89125	0.905	0.876988	0.890492
2	0.92500	0.940	0.927530	0.945016
3	0.97375	0.980	0.973217	0.979408
4	0.99375	0.995	0.993831	0.994876
5	0.99750	0.990	0.997513	0.989837
6	1.00000	0.990	1.000000	0.989837
7	1.00000	0.990	1.000000	0.989837
8	1.00000	0.990	1.000000	0.989837
9	1.00000	0.990	1.000000	0.989837
10	1.00000	0.990	1.000000	0.989837
11	1.00000	0.990	1.000000	0.989837
12	1.00000	0.990	1.000000	0.989837
13	1.00000	0.990	1.000000	0.989837
14	1.00000	0.990	1.000000	0.989837
15	1.00000	0.990	1.000000	0.989837
16	1.00000	0.990	1.000000	0.989837
17	1.00000	0.990	1.000000	0.989837
18	1.00000	0.990	1.000000	0.989837
19	1.00000	0.990	1.000000	0.989837

```
In [134... fig = plt.figure(figsize=(14,10))
fig.subplots_adjust(hspace=0.5, wspace=0.3)

fig.add_subplot(2,2,1)
plt.plot(max_depth, training_accuracy, label='training accuracy')
plt.plot(max_depth, test_accuracy, label='test accuracy')
plt.ylabel('Accuracy',size=20)
plt.xlabel('Max depth',size=20)
plt.xticks(size=15)
plt.yticks(size=15)
plt.title('Accuracy Score',size=20)
plt.legend(['Training','Testing'],frameon=False, ncol=2);

fig.add_subplot(2,2,2)
plt.plot(max_depth, training_f1)
plt.plot(max_depth, test_f1)
plt.ylabel('F1 Score',size=20)
plt.xlabel('Max depth',size=20)
plt.xticks(size=15)
plt.yticks(size=15)
plt.title('F1-Score',size=20)
plt.legend(['Training','Testing'],frameon=False, ncol=2);
```




```
In [135... tree = DecisionTreeClassifier(max_depth=8, min_samples_leaf=1, random_sta
y_pred=tree.predict(X_test)
```


```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.95	0.98	21
1	1.00	0.90	0.95	10
2	0.99	1.00	0.99	169
accuracy			0.99	200
macro avg	1.00	0.95	0.97	200
weighted avg	0.99	0.99	0.99	200

```
In [136... conf_mat = confusion_matrix(y_test,y_pred)
normalized_confusion_matrix(y_test,conf_mat,'Decision Tree')
```



```
In [137... n_features = 13
plt.figure(figsize=(10,6))
plt.barh(range(n_features), tree.feature_importances_, align='center')
plt.yticks(np.arange(n_features), X_train.columns, size=15)
plt.xticks(size=15)
plt.xlabel('Feature importance', size=20)
# plt.ylabel('Feature', size=20)
plt.ylim(-1, n_features)
sns.despine(top=True)
plt.title('Decision Tree', size=20)
plt.tight_layout()
```



We can see that BMI is the most important feature followed by HbA1c.

### Logistic Regression

```
In [138... logreg_cv = LogisticRegression(C=10, solver='lbfgs', max_iter=1000, class
cv_scores = cross_val_score(logreg_cv, X_train_scaled, y_train, cv=5, sco
print("Cross-validation accuracy (5-fold):", cv_scores.mean())

logreg_cv.fit(X_train_scaled, y_train)
y_pred = logreg_cv.predict(X_test)

print(classification_report(y_test, y_pred, zero_division=0))
```

Cross-validation accuracy (5-fold): 0.905

	precision	recall	f1-score	support
0	0.11	1.00	0.19	21
1	0.00	0.00	0.00	10
2	1.00	0.01	0.02	169
accuracy			0.12	200
macro avg	0.37	0.34	0.07	200
weighted avg	0.86	0.12	0.04	200

```
/Users/sushantneopane/Library/Python/3.9/lib/python/site-packages/sklearn/
utils/validation.py:2732: UserWarning: X has feature names, but LogisticRe
gression was fitted without feature names
warnings.warn(
```

```
In [139... import warnings
warnings.filterwarnings("ignore", category=FutureWarning)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

# Use SMOTE to address class imbalance
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

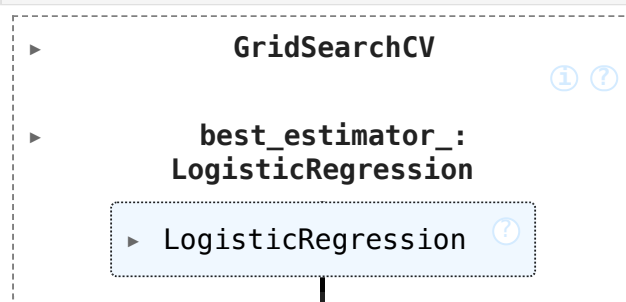
```
# Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_resampled)
X_test_scaled = scaler.transform(X_test)

parameters = {
    'C': [0.001, 0.01, 0.1, 1, 10],
    'penalty': ['l2'],
    'class_weight': [None, 'balanced']
}
```

```
In [140...] logreg_model = LogisticRegression(multi_class='multinomial', solver='lbfgs')

# Perform grid search with cross-validation
logreg_gs = GridSearchCV(logreg_model, parameters, cv=5)
logreg_gs.fit(X_train_scaled, y_train_resampled)
```

Out [140...]



```
In [141...] print(f"Best Parameters: {logreg_gs.best_params_}")

y_pred = logreg_gs.predict(X_test_scaled)
```

Best Parameters: {'C': 10, 'class\_weight': None, 'penalty': 'l2'}


```
In [142...] print('Classification Report (Logistic Regression):\n', classification_re
```

```
Classification Report (Logistic Regression):
```

	precision	recall	f1-score	support
0	0.86	0.86	0.86	21
1	0.43	0.60	0.50	10
2	0.98	0.96	0.97	169
accuracy			0.93	200
macro avg	0.76	0.81	0.78	200
weighted avg	0.94	0.93	0.93	200

```
In [143...] conf_mat = confusion_matrix(y_test, y_pred)
normalized_confusion_matrix(y_test, conf_mat, 'Logistic Regression')
```

## Confusion Matrix (Logistic Regression)



## Clustering

In [144... `print(df)`



	ID	No_Pation	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VL
DL \												
0	502	17975	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	
0.5												
1	735	34221	1	26	4.5	62	4.9	3.7	1.4	1.1	2.1	
0.6												
2	420	47975	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	
0.5												
3	680	87656	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	
0.5												
4	504	34223	1	33	7.1	46	4.9	4.9	1.0	0.8	2.0	
0.4												
..	...	...	...	...	...	..	...	...	...	...	...	
...												
995	200	454317	1	71	11.0	97	7.0	7.5	1.7	1.2	1.8	
0.6												
996	671	876534	1	31	3.0	60	12.3	4.1	2.2	0.7	2.4	1
5.4												
997	669	87654	1	30	7.1	81	6.7	4.1	1.1	1.2	2.4	
8.1												
998	99	24004	1	38	5.8	59	6.7	5.3	2.0	1.6	2.9	1
4.0												
999	248	24054	1	54	5.0	67	6.9	3.8	1.7	1.1	3.0	
0.7												
	BMI	CLASS										
0	24.0	0										
1	23.0	0										
2	24.0	0										
3	24.0	0										
4	21.0	0										
..	...	...										
995	30.0	2										
996	37.2	2										
997	27.4	2										
998	40.5	2										
999	33.0	2										

[1000 rows x 14 columns]

In [145... `df.drop(columns=['ID', 'No_Pation'])`

Out [145...

	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS
0	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0
1	1	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	0
2	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0
3	0	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0
4	1	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	0
...	...	...	...	...	...	...	...	...	...	...	...	...
995	1	71	11.0	97	7.0	7.5	1.7	1.2	1.8	0.6	30.0	2
996	1	31	3.0	60	12.3	4.1	2.2	0.7	2.4	15.4	37.2	2
997	1	30	7.1	81	6.7	4.1	1.1	1.2	2.4	8.1	27.4	2
998	1	38	5.8	59	6.7	5.3	2.0	1.6	2.9	14.0	40.5	2
999	1	54	5.0	67	6.9	3.8	1.7	1.1	3.0	0.7	33.0	2

1000 rows × 12 columns

In [146...

```
y_true = df["CLASS"]
X = df.drop(columns=["CLASS"])
```

In [147...

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

pd.DataFrame(X_scaled, columns=X.columns).head()
```

Out [147...

	ID	No_Pation	Gender	AGE	Urea	Cr	HbA1c	
0	0.672140	-0.074747	-1.139671	-0.401144	-0.144781	-0.382672	-1.334983	-0
1	1.641852	-0.069940	0.877446	-3.130017	-0.212954	-0.115804	-1.334983	-0
2	0.330868	-0.065869	-1.139671	-0.401144	-0.144781	-0.382672	-1.334983	-0
3	1.412950	-0.054126	-1.139671	-0.401144	-0.144781	-0.382672	-1.334983	-0
4	0.680463	-0.069939	0.877446	-2.334096	0.673299	-0.382672	-1.334983	0

In [148...


```
# PCA for visualization
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
```

In [149...

```
inertia = []
K = range(2, 11)
for k in K:
    km = KMeans(n_clusters=k, random_state=42)
    km.fit(X_scaled)
    inertia.append(km.inertia_)

plt.plot(K, inertia, 'bx-')
plt.xlabel('k')
plt.ylabel('Inertia')
```

```
plt.title('Elbow Method for Optimal k (KMeans)')
plt.show()
```




```
In [150... kmeans = KMeans(n_clusters=3, random_state=42)
labels_kmeans = kmeans.fit_predict(X_scaled)

print("KMeans Silhouette Score:", silhouette_score(X_scaled, labels_kmean
print("KMeans Davies-Bouldin Score:", davies_bouldin_score(X_scaled, labe
print("KMeans Calinski-Harabasz Score:", calinski_harabasz_score(X_scaled

plt.scatter(X_pca[:,0], X_pca[:,1], c=labels_kmeans, cmap='viridis', alph
plt.title("KMeans Clusters (PCA Projection)")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.show()
```

KMeans Silhouette Score: 0.10340177510924843  
 KMeans Davies-Bouldin Score: 2.5834372263333667  
 KMeans Calinski-Harabasz Score: 97.61959399467135

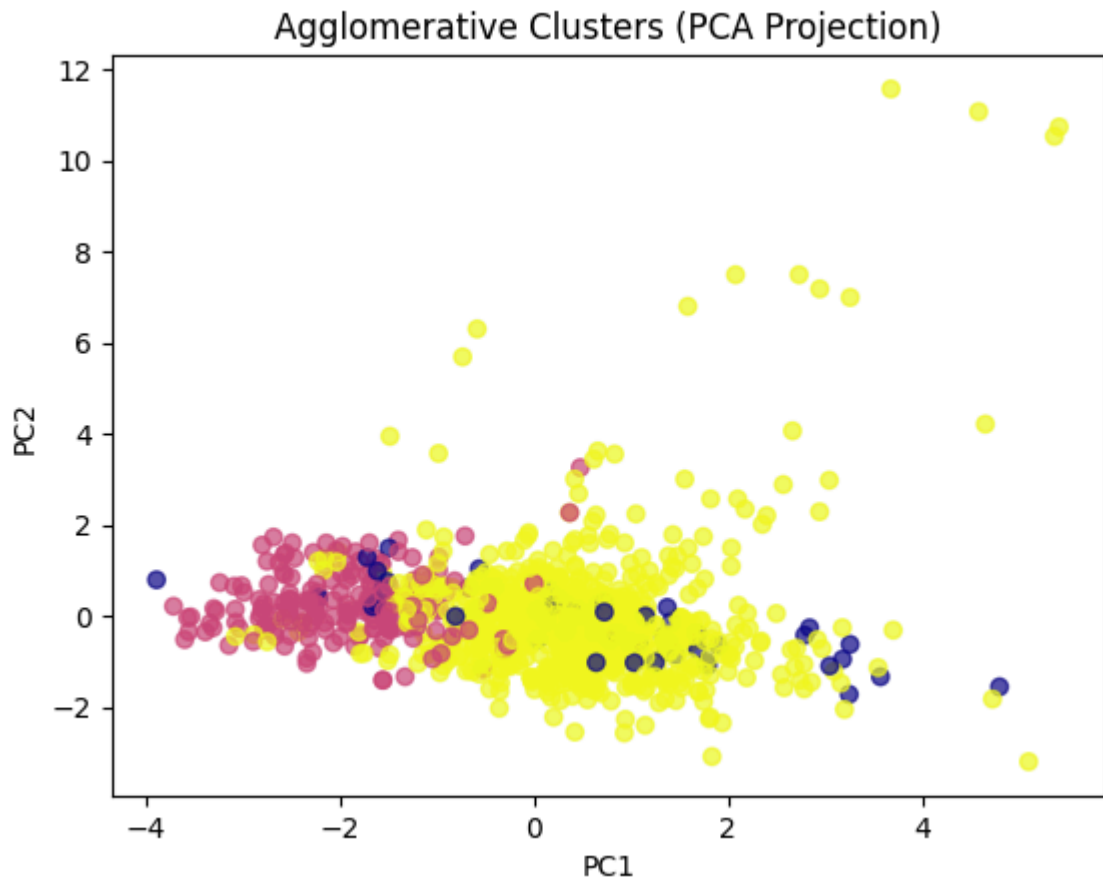


```
In [151... agg = AgglomerativeClustering(n_clusters=3)
labels_agg = agg.fit_predict(X_scaled)

print("Agglomerative Silhouette Score:", silhouette_score(X_scaled, labels_agg))
print("Agglomerative Davies-Bouldin Score:", davies_bouldin_score(X_scaled, labels_agg))
print("Agglomerative Calinski-Harabasz Score:", calinski_harabasz_score(X_scaled, labels_agg))

plt.scatter(X_pca[:,0], X_pca[:,1], c=labels_agg, cmap='plasma', alpha=0.5)
plt.title("Agglomerative Clusters (PCA Projection)")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.show()
```

Agglomerative Silhouette Score: 0.13819106205375598  
 Agglomerative Davies-Bouldin Score: 1.860549541466516  
 Agglomerative Calinski-Harabasz Score: 96.8464672225151



Both K Means Clustering and Agglomerative Clustering perform poorly on the given dataset, based on the clustering evaluation metrics. This indicates that the dataset does not contain strong natural groupings. From the above visualizations, we can also see that there are tons of overlapped classes.

This suggests that unsupervised clustering may not be ideal for this dataset. Given that labels are available, the dataset is likely better suited for supervised learning approaches, where models can learn decision boundaries even in overlapping feature spaces.

## Future Enhancements

- More hyperparameter tuning
- More classifiers such as Random Forest, Gradient Boosting, XGBoost
- Perform feature selection/engineering to reduce noise and redundancy
- Apply dimensionality reduction for better separation and visualization
- Generating a Tableau Dashboard